

**Desalegn Melaku**  
**ECE510**

## **Challenge #12: Accelerating Q-Value Update using PyMTL3 and Cocotb**

### **1. Learning Goals**

This challenge explores the hardware/software co-design of Q-learning components to accelerate reinforcement learning algorithms.

The goals include:

- Estimating execution and communication overhead.
- Deciding on partitioning between hardware and software.
- Selecting rapid prototyping tools like PyMTL3 and cocotb.
- Implementing a full co-simulation with waveform analysis.

### **2. Hardware/Software Co-Design Analysis**

Q-learning agents frequently update their Q-value table using the Bellman equation, which becomes computationally expensive.

We break down the software and hardware overhead as follows:

T1 = Environment step  
T2 = Action selection (e.g.,  $\epsilon$ -greedy)  
T3 = Q-value update in software  
T4 = Logging, tracking  
T5 = Data transfer to hardware  
T6 = Hardware Q-update execution  
T7 = Readback from hardware

For hardware acceleration to be worthwhile:

$$**$T5 + T6 + T7 < T3$**$$

---

Assume:

T3 = 20 ms

T5 = 2 ms

T6 = 4 ms

T7 = 2 ms

Total hardware path = 8 ms < 20 ms → Hardware acceleration justified.

### 3. PyMTL3 Hardware Design: Q-Value Update

The core of the Q-learning update is:

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

This is encoded in PyMTL3 as a combinational update block.

```
class QValueUpdate( Component ):  
    def construct( s ):  
        s.q_old = InPort( Bits32 )  
        s.reward = InPort( Bits32 )  
        s.q_max = InPort( Bits32 )  
        s.alpha = InPort( Bits32 )  
        s.gamma = InPort( Bits32 )  
        s.q_new = OutPort( Bits32 )  
  
        @update  
        def q_update_logic():  
            td = s.reward + (s.gamma * s.q_max) - s.q_old  
            s.q_new @= s.q_old + (s.alpha * td)  
This PyMTL3 model can be translated into Verilog for hardware  
simulation and testing.
```

---

### 4. Cocotb Co-Simulation Testbench

We validate the Q-value update logic using cocotb, a Python-based testbench environment for verifying Verilog modules.

```
@cocotb.test()  
async def test_qvalue(dut):  
    q_old = 10  
    reward = 5  
    q_max = 20  
    alpha = 1  
    gamma = 1
```

```

dut.q_old.value = q_old
dut.reward.value = reward
dut.q_max.value = q_max
dut.alpha.value = alpha
dut.gamma.value = gamma

await RisingEdge(dut.clk)
await Timer(2, units='ns')

expected = q_old + alpha * (reward + gamma * q_max - q_old)
assert dut.q_new.value == expected


```

---

This test confirms functional correctness by comparing the Verilog module output to a Python-calculated expected result.

## 5. Simulation and Waveform Visualization

Below is the simulation waveform. The inputs (Q\_old, reward, Q\_max) are constant over time, while the computed Q\_new value appears in the final cycle.



This confirms the timing of the computation and allows verification of the result propagation through the hardware model.

## 6. Conclusion and Recommendations

This report demonstrated a full co-simulation pipeline for accelerating Q-learning updates:

- PyMTL3 was used for high-level hardware modeling.
- Verilog RTL was generated automatically.
- Cocotb testbench verified correctness and timing.
- A waveform illustrated signal transitions over time.

This co-design approach is ideal for prototyping AI accelerators and testing hardware-software tradeoffs.